

# Demystifying DNS: A Guide to Understanding and Debugging Request Flows in Kubernetes Clusters

Qasim Sarfraz

# About me

- Pakistan (Lahore) -> Germany (Hamburg)
- Software Engineer @ Microsoft
- [Inspektor Gadget](#), [CoreDNS header plugin](#), other cloud-native projects!

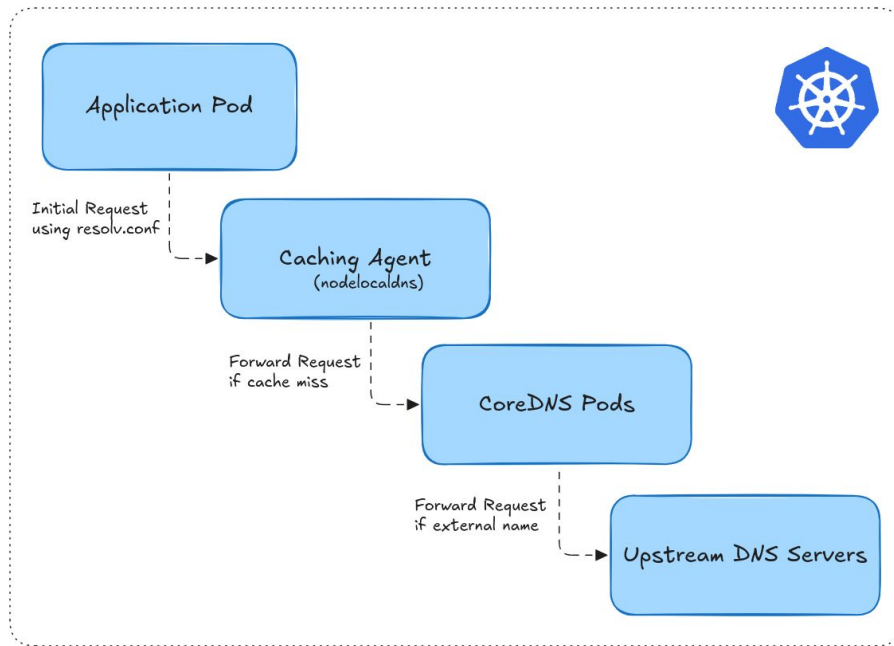


# Agenda

- Understanding DNS Components and Request Flows
- Deep Dive into Request Flows using:
  - CoreDNS log plugin
  - Hubble
  - Inspektor Gadget
- Debugging Scenarios

# Understanding DNS Components / Request Flows

- In order to debug DNS issues we first need to understand different components:
  1. Application Pod
  2. Node Local DNS
  3. kube-dns Service
  4. CoreDNS Pod
  5. Upstream DNS Server
- Challenges:
  - Too many hidden systems.
  - Not easy to trace the DNS request flows across the cluster.
  - Which tools to use to deep dive into request flows?



# Deep Dive into Request Flows

# Deep Dive into Request Flows - Disclaimer

We will trace a DNS request (name=**example.com.**) from pod (name=**mypod**)  
(namespace=**demo**) on a **minikube** Kubernetes Cluster

# CoreDNS log plugin

- CoreDNS core plugin: <https://coredns.io/plugins/log/>
- Needs a keyword **log** in Corefile.
- Logs all the requests to stdout.

```
log [NAMES...] [FORMAT] {  
    class CLASSES...  
}
```

# CoreDNS log plugin

```
$ kubectl logs -n kube-system -l k8s-app=kube-dns -f  
[INFO] 10.244.0.173:34432 - 45551 "A IN example.com. udp 29 false 512" NOERROR qr,rd,ra 56 0.021960059s
```



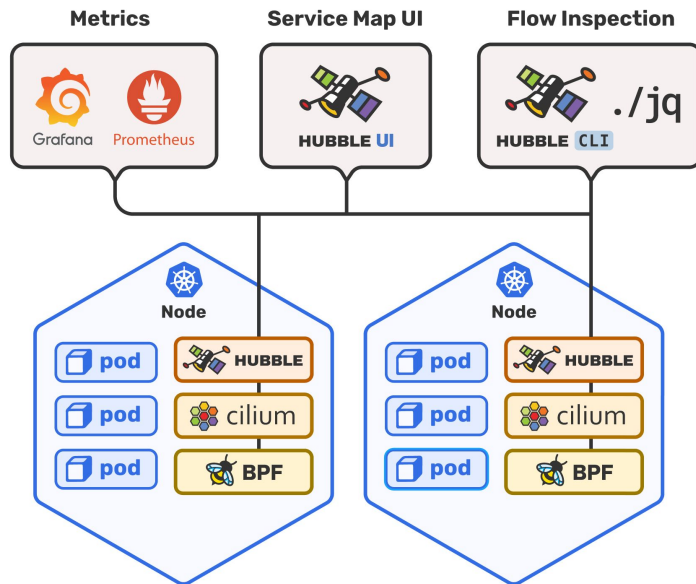
Log Format:

```
{remote}:{port} - {>id} "{type} {class} {name} {proto} {size} {>do} {>bufsize}" {rcode} {>rflags} {rsize} {duration}
```



# Hubble

- Provides Network, Service & Security Observability for Kubernetes.
- Built on top of Cilium and eBPF.
- Installed as Kubernetes Daemonset.



source: <https://github.com/cilium/hubble>

# Hubble

- Install Cilium with L7 Proxy Support.
- Create cilium network policies for visibility of DNS traffic for:
  - mypod
  - CoreDNS

```

1  apiVersion: cilium.io/v2
2  kind: CiliumNetworkPolicy
3  metadata:
4    name: mypod-visibility
5    namespace: demo
6  spec:
7    endpointSelector:
8      matchLabels:
9        run: mypod
10   egress:
11     - toEntities:
12       - all
13     - toEndpoints:
14       - matchLabels:
15         k8s:io.kubernetes.pod.namespace: kube-system
16         k8s:k8s-app: kube-dns
17     toPorts:
18       - ports:
19         - port: "53"
20         protocol: ANY
21     rules:
22       dns:
23         - matchPattern: '*'

```

mypod

```

1  apiVersion: cilium.io/v2
2  kind: CiliumNetworkPolicy
3  metadata:
4    name: kube-dns-visibility
5    namespace: kube-system
6  spec:
7    endpointSelector:
8      matchLabels:
9        k8s-app: kube-dns
10   egress:
11     - toEntities:
12       - all
13     - toPorts:
14       - ports:
15         - port: "53"
16         protocol: ANY
17     rules:
18       dns:
19         - matchPattern: '*'

```

CoreDNS

```
$ hubble observe --protocol dns -f
```

```
Aug 26 13:03:59.569: demo/mypod:49105 (ID:39877) -> kube-system/coredns-6948557899-m7d49:53 (ID:22818) dns-request proxy FORWARDED (DNS Query example.com. A)
Aug 26 13:03:59.570: kube-system/coredns-6948557899-m7d49:43347 (ID:22818) -> 192.168.49.1:53 (world) dns-request proxy FORWARDED (DNS Query example.com. A)
Aug 26 13:03:59.572: kube-system/coredns-6948557899-m7d49:43347 (ID:22818) <- 192.168.49.1:53 (world) dns-response proxy FORWARDED (DNS Answer "93.184.215.14" TTL: 3091 (Proxy example.com. A))
Aug 26 13:03:59.572: demo/mypod:49105 (ID:39877) <- kube-system/coredns-6948557899-m7d49:53 (ID:22818) dns-response proxy FORWARDED (DNS Answer "93.184.215.14" TTL: 30 (Proxy example.com. A))
```

Filtering:

```
$ nslookup -query=a unknown.example.com.
```

```
$ hubble observe --since=5m --protocol dns -o json \
```

```
| jq 'select(.flow.l7.dns.rcode==3) | .flow.destination.namespace + "/" + .flow.destination.pod_name' \
```


```
| sort | uniq -c | sort -r
```

```
1 "kube-system/coredns-6948557899-m7d49"
```

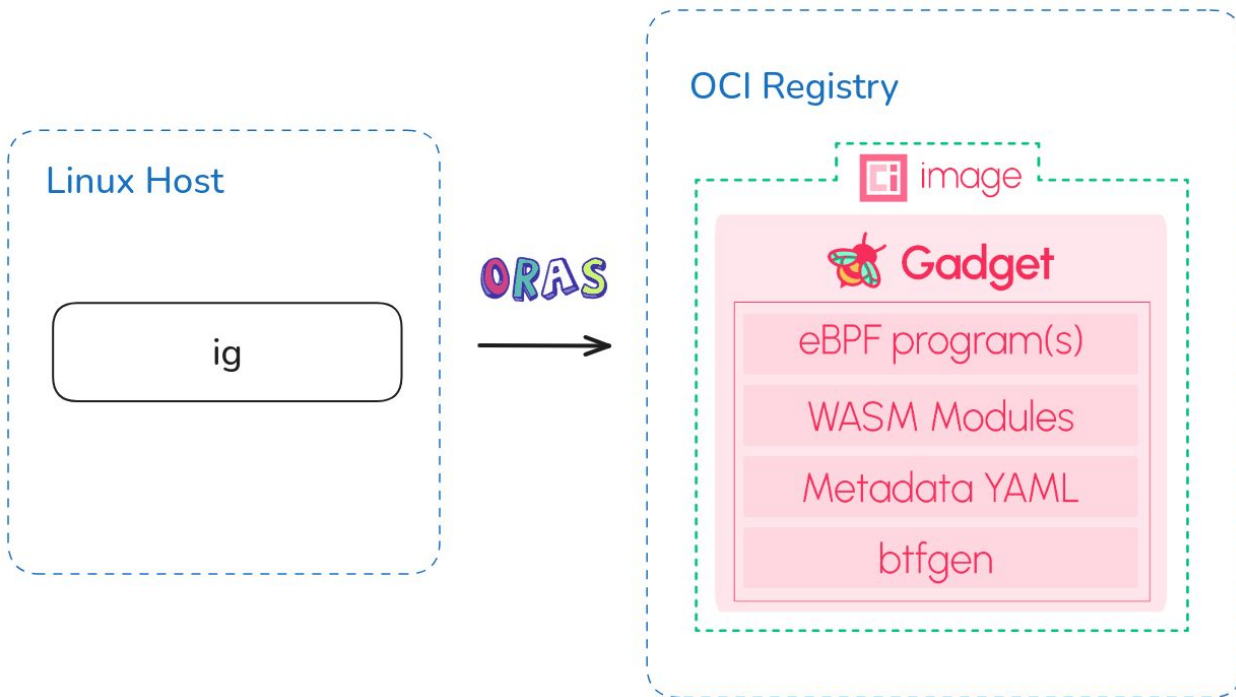
```
1 "demo/mypod"
```

# Inspektor Gadget

- **Tool:** A set of tools (gadgets) that empower users to inspect Kubernetes and Linux Systems.
- **Framework:** An observability framework that allows you to build and share custom “gadgets”.
- It can be used in Kubernetes as:
  - Daemonset + kubectl gadget.
  - kubectl debug node.



# Inspektor Gadget - Gadget



# Inspektor Gadget - Gadget

The screenshot displays the Artifact HUB web interface. At the top, the 'ArtifactHUB' logo is on the left, a search bar with the placeholder 'Search packages' is in the center, and navigation links for 'DOCS', 'STATS', and a user profile icon are on the right. Below the header, it indicates '1 - 18 of 18 results'. A filter bar shows 'Filters: Official, Verified publisher, CNCF, KIND: Inspektor gadgets'. On the left side, a 'FILTERS' sidebar includes a 'Reset' button and checkboxes for 'Official', 'Verified publishers', and 'CNCF', all of which are checked. Under the 'KIND' section, 'Inspektor gadgets' is selected. The main content area shows two gadget cards. The first card is for 'snapshot process', an 'Inspektor gadget' that is 'Official', updated 20 days ago, and at version 0.31.0. It features a red shield icon and a 'Monitoring and logging' tag. The second card is for 'snapshot socket', also an 'Inspektor gadget' that is 'Official', updated 20 days ago, and at version 0.31.0, featuring a red shield icon. Both cards include a 'Show' button and a set of four small icons representing different functionalities.

# Inspektor Gadget - DNS Gadget

- [DNS Gadget](#).
- Traces DNS requests and responses using eBPF.
- Uses WASM for post-processing.
- [Source code](#). (Modify -> Build -> Push -> Run)



Guide



Source Code

# Inspektor Gadget (Application Pod)

```
$ kubectl gadget run trace_dns:latest -n demo -p mypod
```

K8S.NODE	K8S.NAMESPACE	K8S.PODNAME	K8S.CONTAINERNAME	SRC	DST	COMM	PID	QR	QTYPE	NAME	RCODE	ADDRESSES
minikube	demo	mypod	mypod	10.244.0.4:48899	10.96.0.10:53	isc-net-0...	1392643	Q	A	example.com.		
minikube	demo	mypod	mypod	10.96.0.10:53	10.244.0.4:48899	isc-net-0...	1392643	R	A	example.com.	Success	93.184.2...

Kubernetes Enrichment

mypod IP

kube-dns Service IP

Query Information

Response Code



# Inspektor Gadget (Application + CoreDNS Pods)

```
$ kubectl gadget run trace_dns:latest -n demo,kube-system -F "k8s.podName~mypod|coredns-.*" -F "name==example.com."
--fields=k8s.node,k8s.namespace,k8s.podname,id,src,dst,qr,name,rcode,address
```

K8S.NAMESPACE	K8S.PODNAME	ID	SRC	DST	QR	QTYPE	NAME	RCODE	ADDRESSES
demo	mypod	a9ba	10.244.0.4:44601	10.96.0.10:53	Q	A	example.com.		
kube-system	coredns-7db6d8ff4d-nlmmg	a9ba	10.244.0.4:44601	10.244.0.2:53	Q	A	example.com.		
kube-system	coredns-7db6d8ff4d-nlmmg	2c39	192.168.49.1:53	10.244.0.2:45117	R	A	example.com.	Success	93.184.215.14
kube-system	coredns-7db6d8ff4d-nlmmg	2c39	10.244.0.2:45117	192.168.49.1:53	Q	A	example.com.		
kube-system	coredns-7db6d8ff4d-nlmmg	a9ba	10.244.0.2:53	10.244.0.4:44601	R	A	example.com.	Success	93.184.215.14
demo	mypod	a9ba	10.96.0.10:53	10.244.0.4:44601	R	A	example.com.	Success	93.184.215.14

# Inspektor Gadget (complete trace with OS context)

```
$ kubectl gadget run trace_dns:latest -n demo,kube-system -F "name==example.com."
```

```
--fields=k8s.namespace,k8s.podname,id,src,dst,qr,qtype,name,rcode,addresses,pkt_type,netns_id,latency_ns,timestamp
```

K8S.NAMESPACE	K8S.PODNAME	ID	SRC	DST	QR	QTYPE	NAME	RCODE	ADDRESSES	PKT_TY...	NETNS_ID	LATENC...	TIMESTAMP
demo	mypod	c1a1	10.244.0.4:50202	10.96.0.10:53	Q	A	example.com.			OUTGOI...	4026532572	0	2024-08-26T15:45:49.556567446Z
		c1a1	10.244.0.4:50202	10.96.0.10:53	Q	A	example.com.			OTHERH...	4026532220	0	2024-08-26T15:45:49.556583162Z
		c1a1	10.244.0.4:50202	10.244.0.2:53	Q	A	example.com.			OUTGOI...	4026532220	0	2024-08-26T15:45:49.556649793Z
kube-system	coredns-7db6d8ff4d-nlmmg	c1a1	10.244.0.4:50202	10.244.0.2:53	Q	A	example.com.			HOST	4026532859	0	2024-08-26T15:45:49.556658132Z
kube-system	coredns-7db6d8ff4d-nlmmg	0aa5	10.244.0.2:38844	192.168.49.1:53	Q	A	example.com.			OUTGOI...	4026532859	0	2024-08-26T15:45:49.557054752Z
		0aa5	10.244.0.2:38844	192.168.49.1:53	Q	A	example.com.			OTHERH...	4026532220	0	2024-08-26T15:45:49.557066317Z
		0aa5	192.168.49.1:53	10.244.0.2:38844	R	A	example.com.	Success	93.184.215.14	OUTGOI...	4026532220	0	2024-08-26T15:45:49.557978226Z
		0aa5	192.168.49.1:53	10.244.0.2:38844	R	A	example.com.	Success	93.184.215.14	OUTGOI...	4026532220	0	2024-08-26T15:45:49.557990376Z
kube-system	coredns-7db6d8ff4d-nlmmg	0aa5	192.168.49.1:53	10.244.0.2:38844	R	A	example.com.	Success	93.184.215.14	HOST	4026532859	942911	2024-08-26T15:45:49.557997663Z
kube-system	coredns-7db6d8ff4d-nlmmg	c1a1	10.244.0.2:53	10.244.0.4:50202	R	A	example.com.	Success	93.184.215.14	OUTGOI...	4026532859	0	2024-08-26T15:45:49.558255874Z
		c1a1	10.244.0.2:53	10.244.0.4:50202	R	A	example.com.	Success	93.184.215.14	OTHERH...	4026532220	0	2024-08-26T15:45:49.558264967Z
		c1a1	10.96.0.10:53	10.244.0.4:50202	R	A	example.com.	Success	93.184.215.14	OUTGOI...	4026532220	0	2024-08-26T15:45:49.558285510Z
demo	mypod	c1a1	10.96.0.10:53	10.244.0.4:50202	R	A	example.com.	Success	93.184.215.14	HOST	4026532572	1723720	2024-08-26T15:45:49.558291166Z



Initial Client DNS Request being forwarded via HOST



Upstream DNS Request/Response being forwarded via HOST



Initial Client DNS Response being forwarded via HOST

# Deep Dive into Request Flows - Summary

CoreDNS log plugin	Hubble	Inspektor Gadget
Ideal for initial inspection; however, its scope is limited to CoreDNS.	Provides a compact overview of request flows with enrichment and retrospective analysis.	Offers rich DNS traces with OS context and Kubernetes enrichment/filtering.
Requires configuration changes for use.	Requires Cilium CNI / specific policies to be in place for L7 flow visibility.	Extensible through custom gadget images ; No Support for TCP.

# Debugging Scenarios

# Debugging Scenario 1 - Verify Health of an Upstream DNS Server

## Debugging Scenario 2 - Identify Unsuccessful DNS Response

# Conclusion

- First step for debugging DNS is understanding the request flows and:
  - [CoreDNS log plugin](#).
  - [Hubble](#).
  - [Inspektor Gadget](#).
- This was just a starting point so feedback is appreciated.
- Finally, reach out communities if you have any questions/suggestions.

# Thank you for Listening!