



Debugging Kubernetes and Container Workloads with the Superpower of eBPF

Kubernetes Community Day, Islamabad, 2024



Hello!



Qasim Sarfraz

Software Engineer @ Microsoft

- I am from Lahore, Pakistan
- Currently, based in Hamburg, Germany
- I work with on OSS, Containers, Kubernetes
- I'm focused on [Inspektor Gadget](#), [kubectl aks](#) and [CoreDNS header plugin](#).
- I'm available at <https://mqasimsarfraz.com>

Agenda

1. What is eBPF?
2. Debugging Containers using eBPF
3. Inspektor Gadget as a **tool**
4. Inspektor Gadget as a **framework**
5. Contributing

What is eBPF?

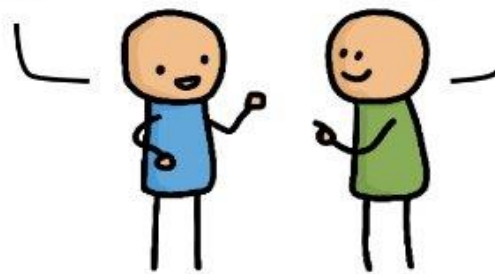
Application Developer:

I want this new feature to observe my app



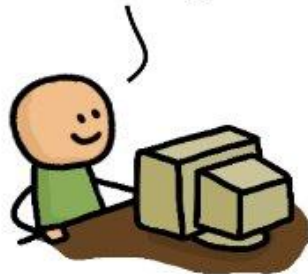
Hey kernel developer! Please add this new feature to the Linux kernel

OK! Just give me a year to convince the entire community that this is good for everyone.



1 year later...

I'm done. The upstream kernel now supports this.



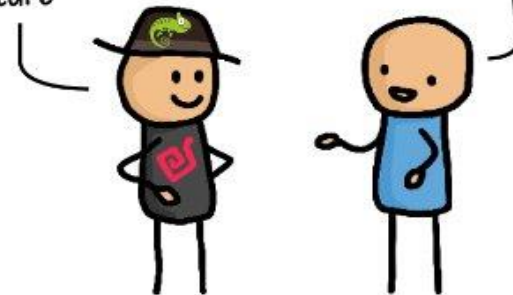
But I need this in my Linux distro



5 year later...

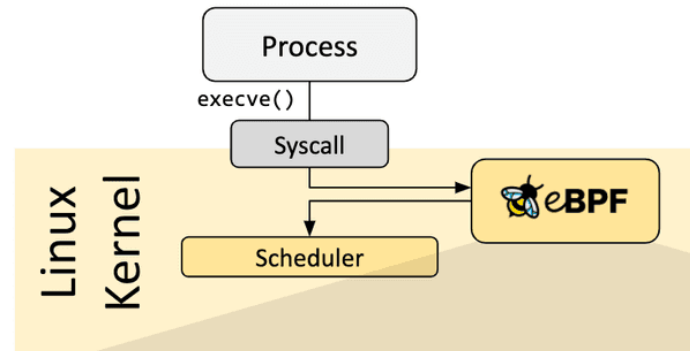
Good news. Our Linux distribution now ships a kernel with your required feature

OK but my requirements have changed since...



What is eBPF?

“eBPF does to Linux what JavaScript does to HTML. (Sort of.)” - Brendan Gregg, Netflix



```
int syscall__ret_execve(struct pt_regs *ctx)
{
    struct comm_event event = {
        .pid = bpf_get_current_pid_tgid() >> 32,
        .type = TYPE_RETURN,
    };

    bpf_get_current_comm(&event.comm, sizeof(event.comm));
    comm_events.perf_submit(ctx, &event, sizeof(event));

    return 0;
}
```

Debugging Containers using eBPF

Debugging Containers using eBPF

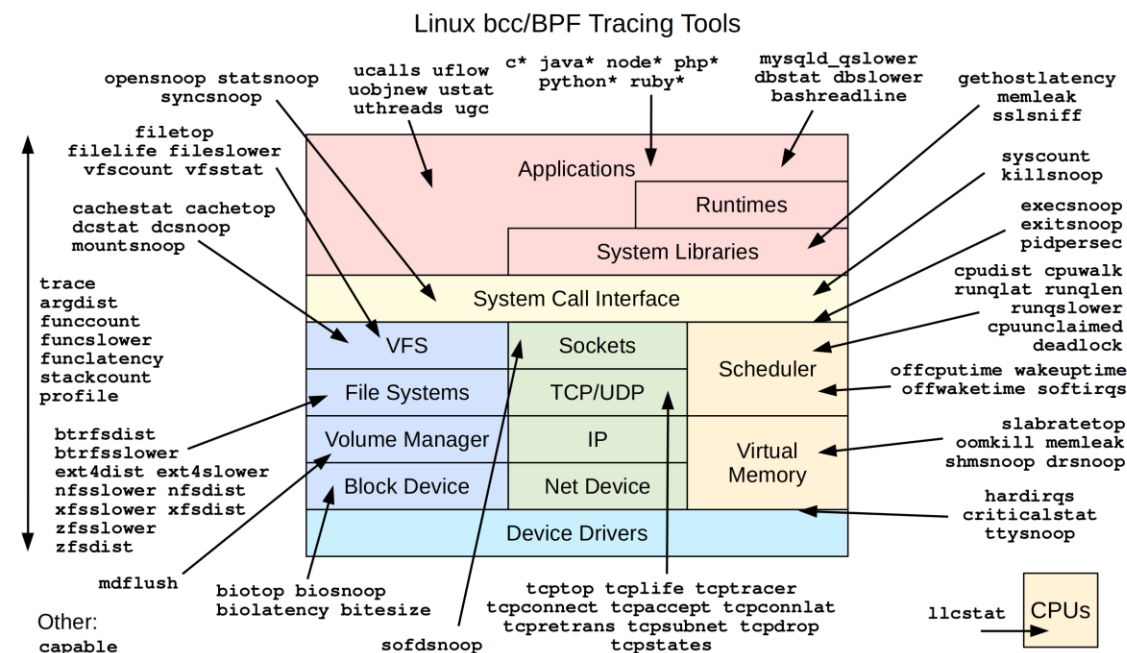
- BPF Compiler Collection (BCC):

- Special Filtering:

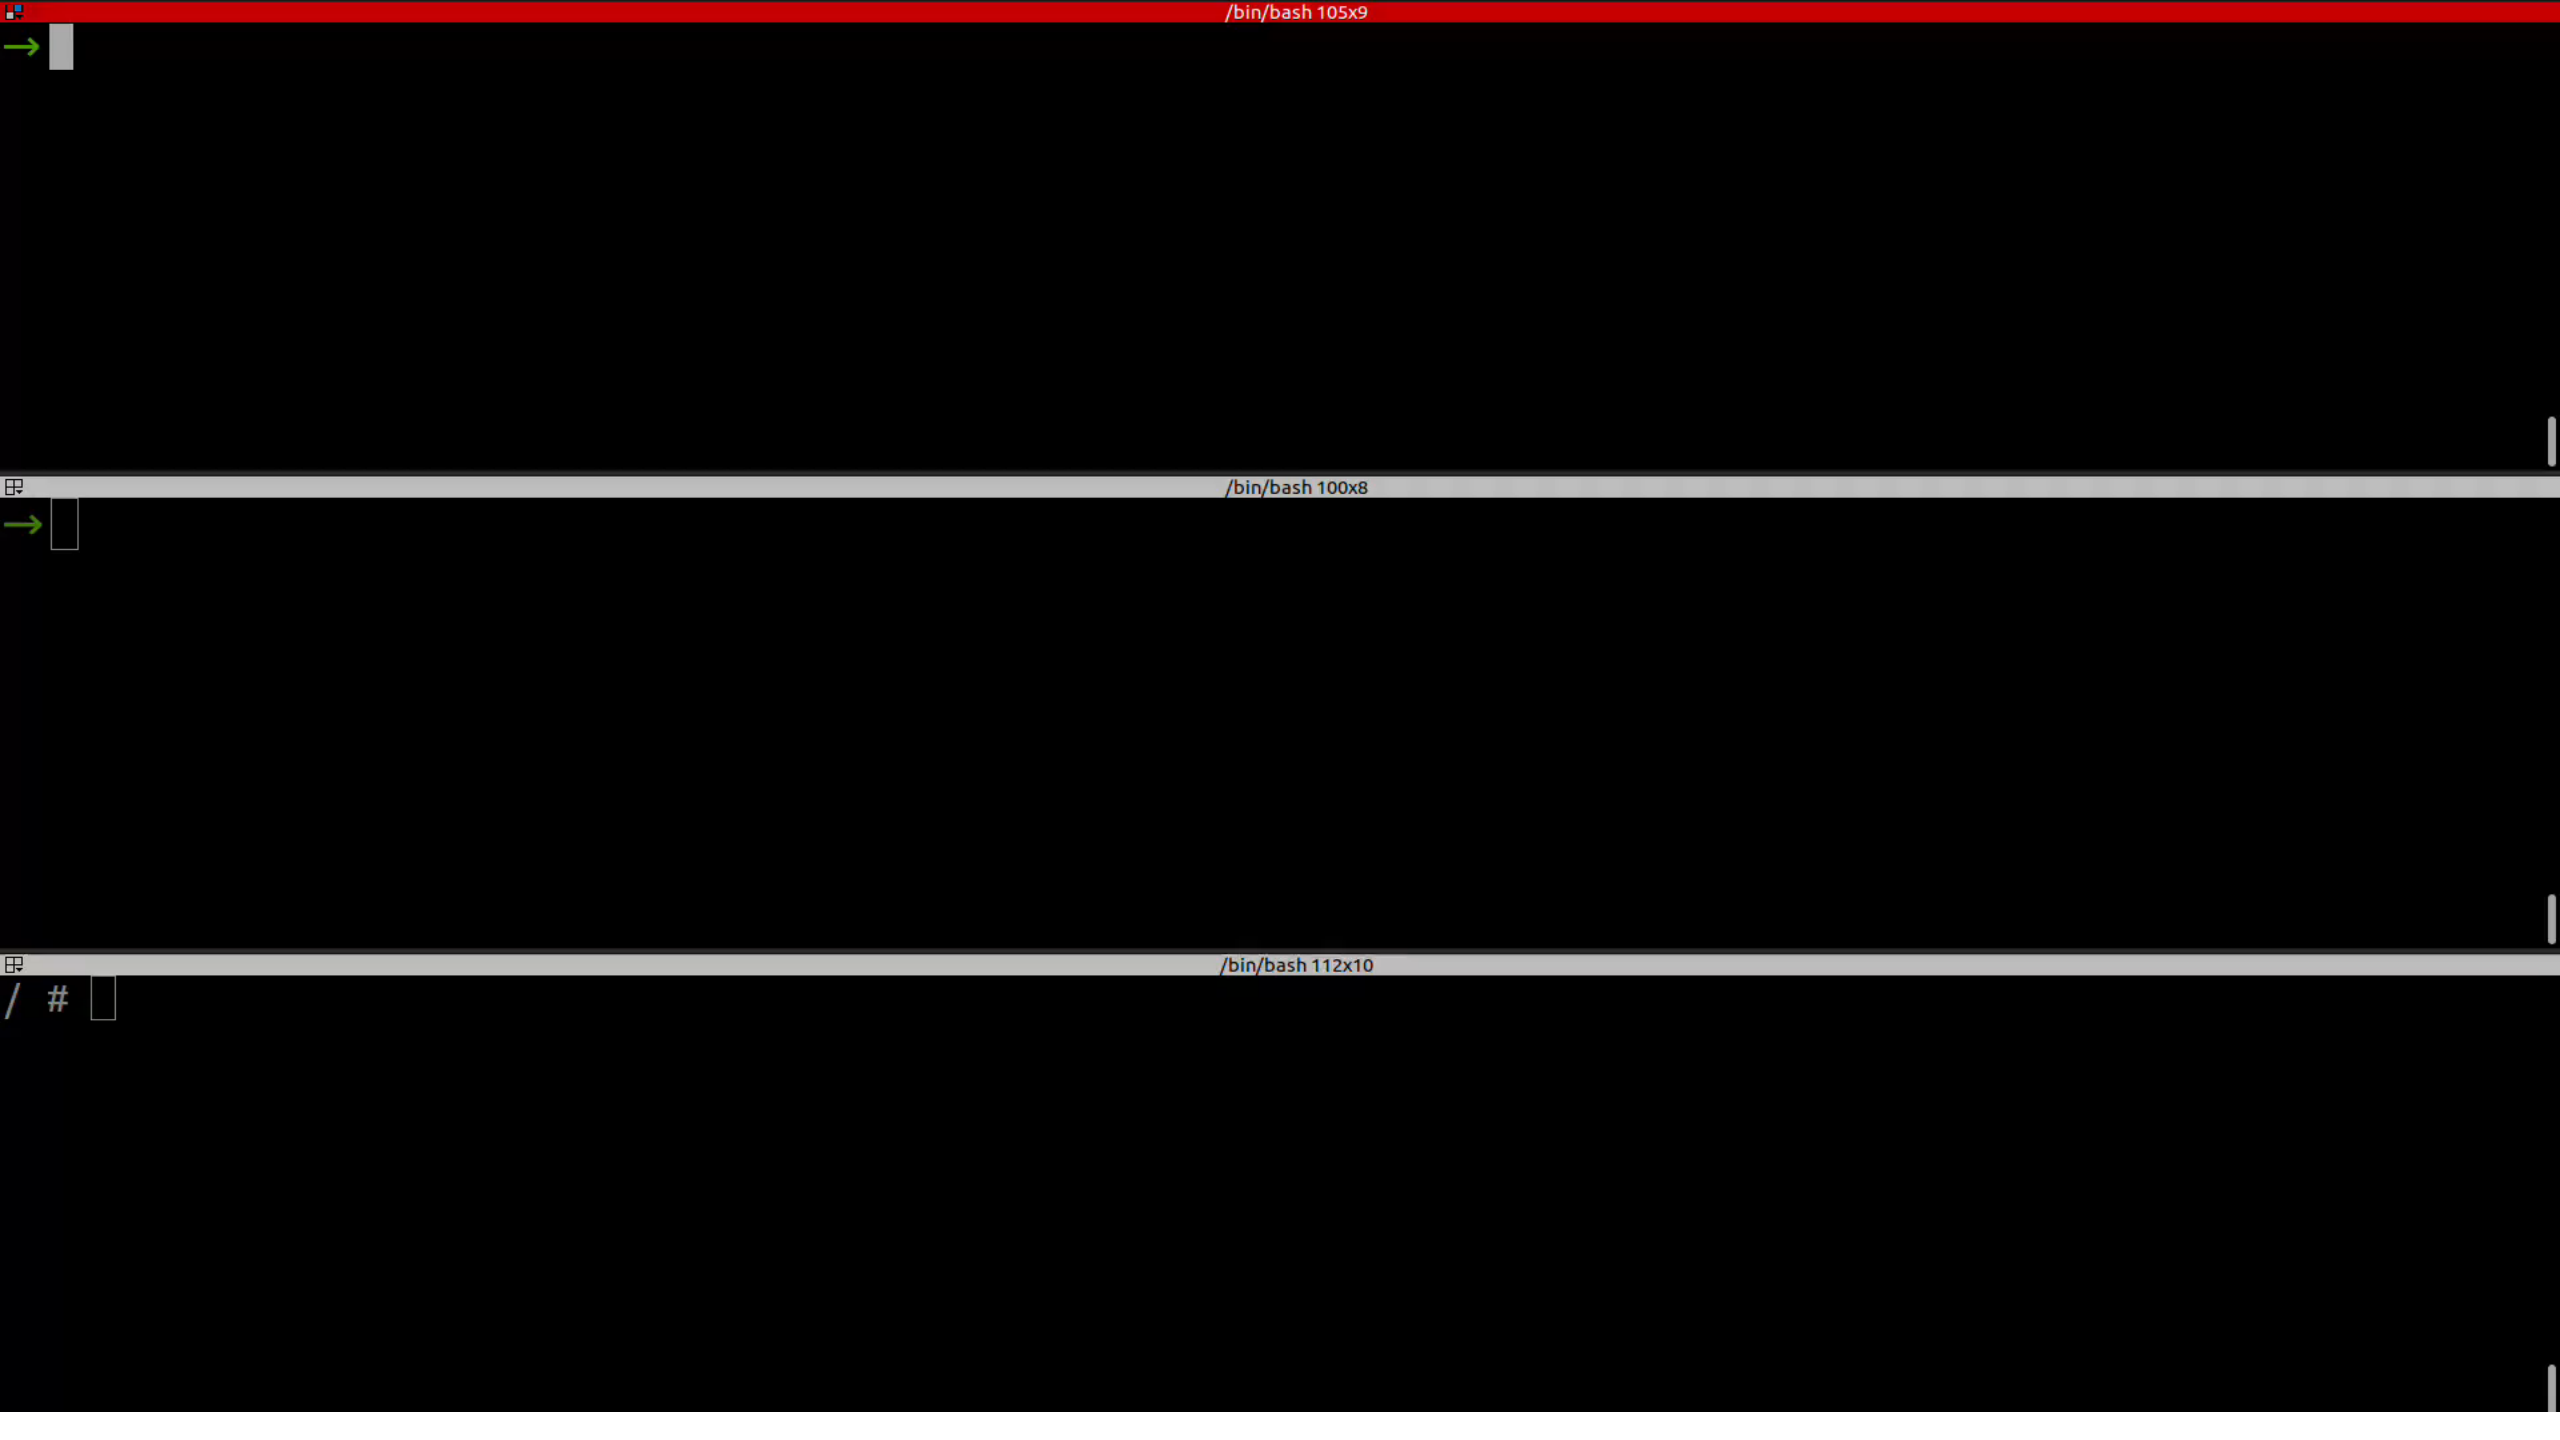
- By cgroups
- By mount namespace

- For Example:

- Create an an eBPF map
- Start a BCC tool with the path to eBPF map
- Update eBPF map by adding mount namespace of the container



<https://github.com/iovisor/bcc#tools> 2019



- Painful debugging experience.
- Lack of support of:
 - Container enrichment
 - Container filtering
- Distribution/Packaging Problem?

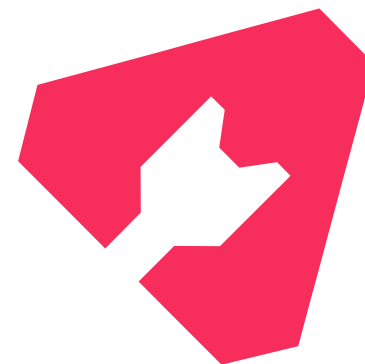
Inspektor Gadget as a tool

<https://www.inspektor-gadget.io>

Inspektor Gadget as a tool



- A collection of **eBPF-based** gadgets to debug and inspect Kubernetes apps and resources
- It is a **CNCF** sandbox project
- Available as a krew plugin
 - `kubectl krew install gadget`
 - `kubectl gadget deploy`
- Also available as:
 - A CLI tool for Linux hosts called **ig**
 - A container image for `kubectl debug ...`



**INSPEKTOR
GADGET**



Can we improve it?

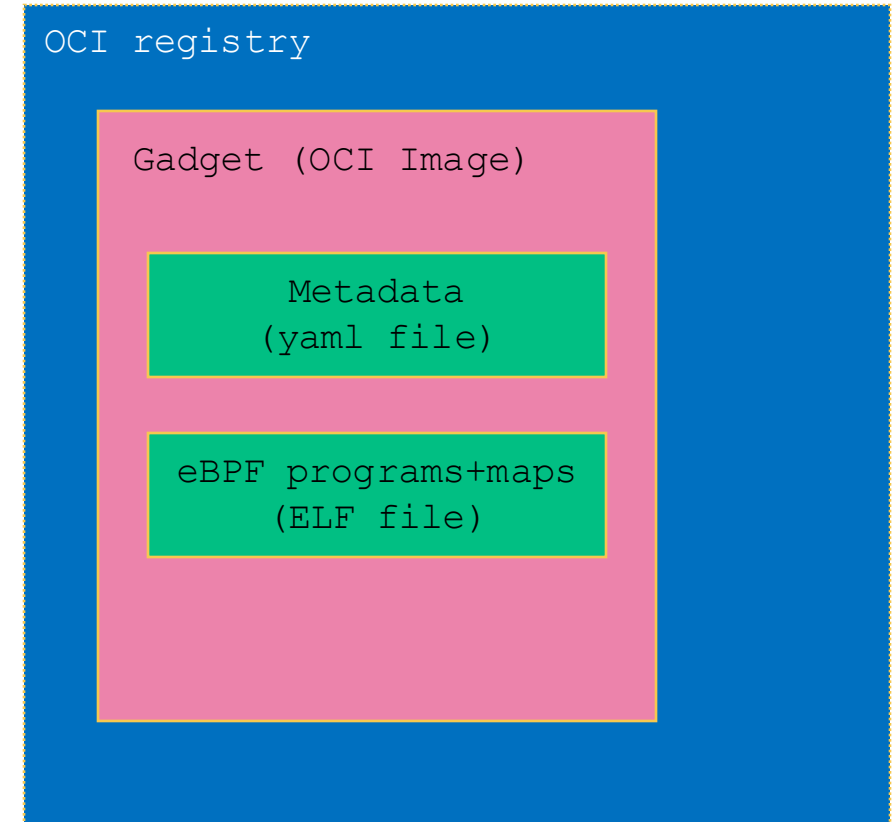
- Painful debugging experience. ✓
- Lack of support of:
 - Container enrichment ✓
 - Container filtering
- Distribution/Packaging Problem ?

Inspektor Gadget as a framework

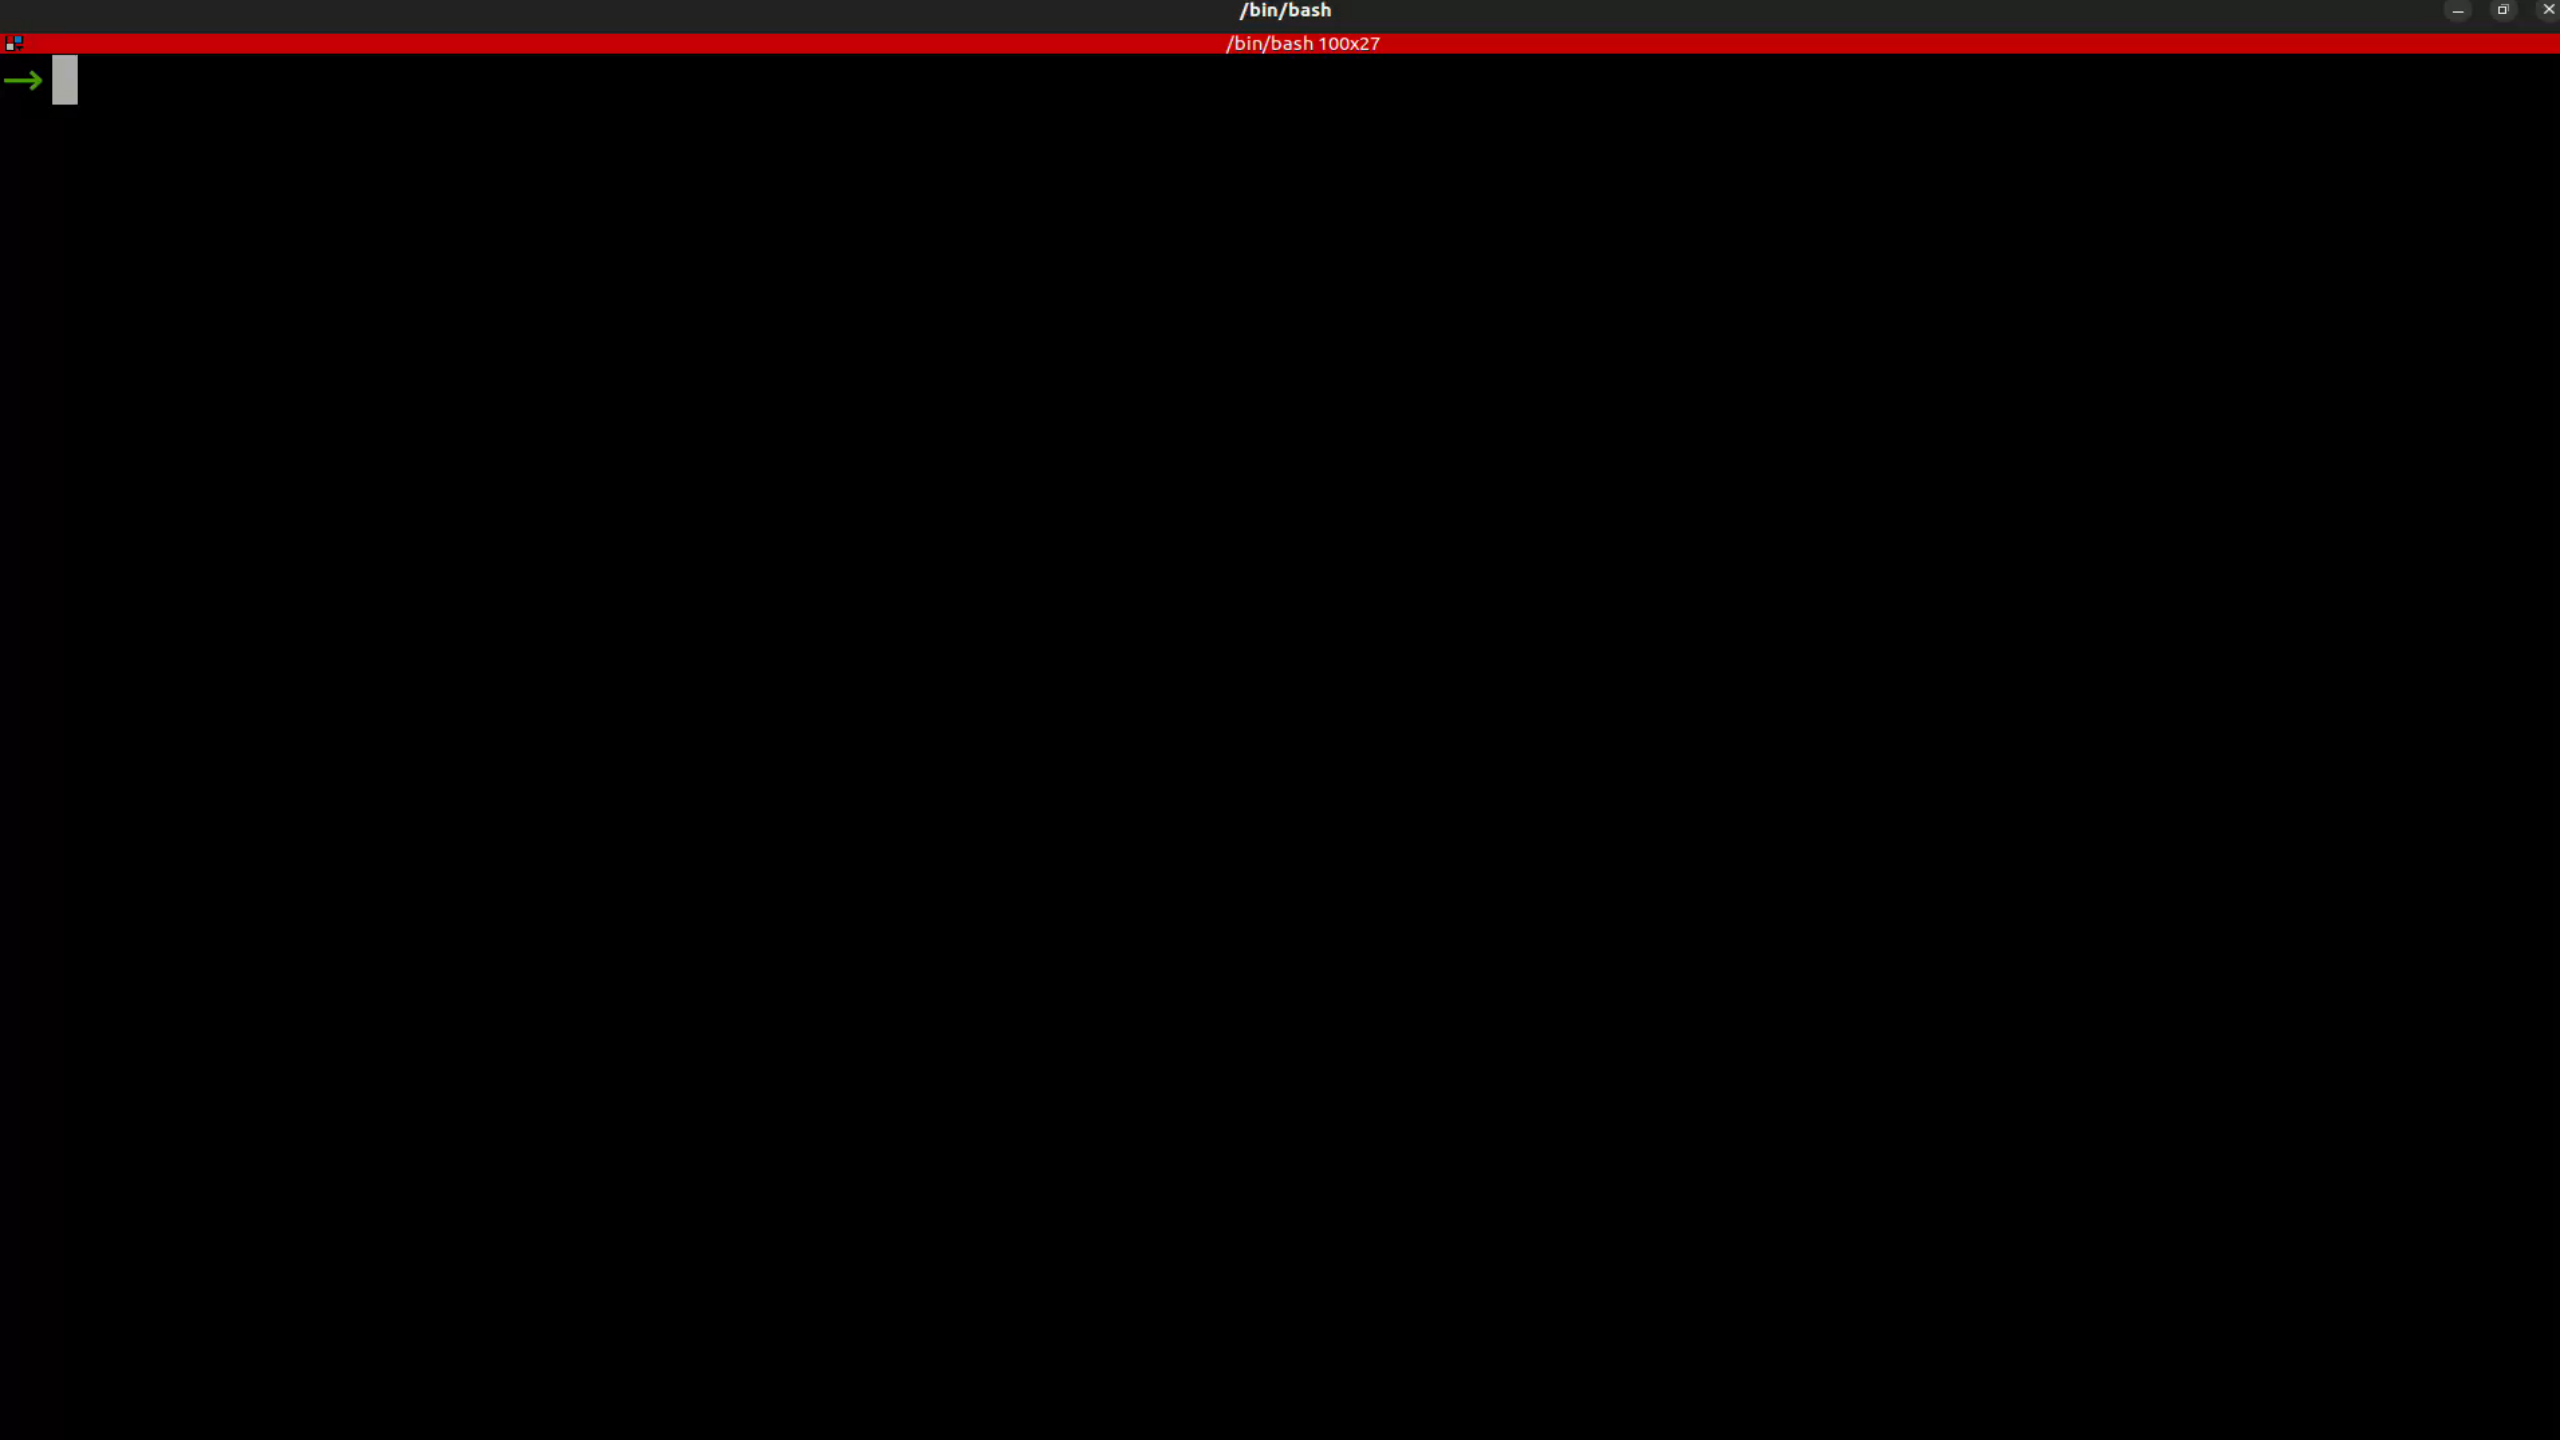
<https://www.inspektor-gadget.io>

Inspektor Gadget as a framework

- Decouple the idea of gadgets from Inspektor Gadget
- Use the idea of [OCI artifacts](#) e.g. SBOMS, Helm Charts etc.
- Gadgets packaged as OCI images and Published to OCI registries



- A familiar UX to package your eBPF programs:
 - **build** – builds a gadget as an OCI image
 - **push** – pushes the gadget (OCI image) to a registry
 - **pull** – pulls the gadget from a registry
 - **run** – runs the specified gadget
- Discovery via [Artifacthub](https://artifacthub.io)
- The feature is under active development so looking for feedback from community! :)



Interested in Contributing?

- We live on [Github](#)
- Use [CONTRIBUTING.md](#) as a starting point
- Look for issues with label: ["good first issue"](#)
- Reach us out at Kubernetes slack: [#inspektor-gadget](#)



Thank you